

## Section 5.2 Processes

1. Overview
2. Process management
3. Inter-process communications

### 5.2.1 Overview

- ◆ What is a process?
  - a running executable
- ◆ Management of processes
  - by program user
    - ✱ using shell commands
  - by other programs
    - ✱ using system calls

## Overview (cont.)

- ◆ Each process has:
  - unique process identifier (PID)
  - parent process (PPID)
    - ✱ the process that spawned it
  - address space and virtual memory
    - ✱ code segment, data segment, function call stack, heap
  - control flow(s)

## 5.2.2 Process Management

- ◆ Processes are typically managed by OS
- ◆ OS also allows users to manage processes
  - from shell command line
  - from another program

# Process Management From Shell

## ◆ From a shell, a user can:

- start process
  - \* in foreground
  - \* in background
- send signal to process
  - \* suspend
  - \* stop
  - \* ... more on this later ...

# Process Management and System Calls

## ◆ A program can:

- start a new process by cloning itself
  - \* **fork** system call
- start a new process by morphing itself
  - \* **exec** family of system calls


# Forking a Clone Process

`pid_t fork(void)`

## ◆ Description:

- creates a clone of the current process
  - \* current process is the *parent*
  - \* new process is the *child*
  - \* child process gets copy of parent's address space
- return value
  - \* in child process
    - ◆ zero
  - \* in parent process
    - ◆ child process id if successful
    - ◆ -1 in case of error

# Forking a Clone Process (cont.)

- ◆ Multiple child processes can be spawned
  - child processes get a copy of parent code
  - multiple ~~forks~~ in the parent mean multiple ~~forks~~ in the children
- ◆ Watch for fork bombs 
  - OS keeps process table
  - all tables have finite capacity

# Morphing Into Another Process

## ♦ **exec** family of system calls

- replace executing code of current process with another program
  - \* same PID
  - \* different instructions
- include `execl()`, `execle()`, `execlp()`, `execvp()`, `execv()`
- differences in parameters and environment settings
- if **exec** call fails, original program continues

# Waiting for a Child Process

```
pid_t wait(int *status)
```

## ♦ Description:

- pauses execution of parent until any child process terminates
- return value
  - \* child pid if successful
  - \* -1 in case of error

## Waiting for a Child Process (cont.)

```
pid_t waitpid(pid_t pid, int *status, int options)
```

### ◆ Description:

- pauses execution of parent until specified child process terminates
- return value
  - \* child pid if successful
  - \* -1 in case of error

## Invoking a Shell Command

```
int system(const char *command)
```

### ◆ Description:

- runs the specified `command` as a shell command
- process blocks until command execution has completed
- return value
  - \* shell process status if successful
  - \* -1 in case of error

## 5.2.3 Inter-Process Communications

- ◆ IPC overview
- ◆ Signals
- ◆ Sockets

## IPC Overview

- ◆ What is inter-process communications (IPC)?
  - sending and receiving information between processes
    - ✱ on the same physical host
    - ✱ on separate physical hosts
      - ◆ must be networked
- ◆ Main approaches to IPC
  - signals
  - sockets

# Signals

## ◆ What is a signal?

- a value (integer) sent from one process to another
  - ✳ there is a fixed set of existing signal values (30 to 40)
    - ◆ `/usr/include/.../bits/signal.h`
  - ✳ only two are user-defined
  - ✳ can be sent from shell too!
- typically used in error situations
  - ✳ tell program to terminate
- very limited kind of IPC
  - ✳ processes must be on the **same** host

# Signals (cont.)

## ◆ Two steps in using signals

- install a signal handler
  - ✳ indicate which function is called when a specific signal is received
- send a signal
  - ✳ send a specific signal from one process to another



# Installing a Signal Handler

- ◆ What is a signal handler?
  - a function called when a specific signal is received
- ◆ Characteristics
  - every signal has its own handler
  - there is a default handler for every signal
    - \* usually terminates the program
  - signal handler is installed using `signal` system call

# Installing a Signal Handler (cont.)

```
sighandler_t signal(int signum, sighandler_t action)
```

- ◆ Description:
  - installs signal handler specified in `action` to handle signal `signum`
  - `sighandler_t` is a predefined type
    - \* used for function that takes one `int` as parameter and returns `void`
  - returns signal handler previously associated with `signum`

## Installing a Signal Handler (cont.)

### ◆ Description (cont.):

- `signum` must be one of the predefined signal values
- `action` can have one of the following values:
  - \* `SIG_IGN`
    - ◆ tells OS to ignore the signal and do nothing
  - \* `SIG_DFL`
    - ◆ tells OS to call the default signal handler
  - \* a signal handler function
    - ◆ tells OS to call the specified function

## Sending a Signal

```
int kill(pid_t pid, int signum)
```


### ◆ Description:

- sends the signal `signum` to the process with identifier `pid`
- `signum` must be one of the predefined signal values
- return value
  - \* 0 if successful
  - \* -1 in case of error

# Sockets

- ◆ What is a socket?
  - an endpoint in IPC
    - ✱ processes can be on same or different hosts
  - socket address made up of:
    - ✱ IP address
      - ◆ indicating a unique host
    - ✱ port number
      - ◆ indicating a unique application running on that host
  - represented as an integer

# Basic Networking



- ◆ Network layer protocol:
  - Internet Protocol (IP)
- ◆ Transport layer protocols:
  - Transmission Control Protocol (TCP)
  - User Datagram Protocol (UDP)

# Socket Components

## ◆ IP address

- uniquely identifies a computer at the network layer

## ◆ Port number

- uniquely identifies a process at the transport layer
- only specific range of values is unreserved

# Types of Sockets

## ◆ Stream sockets

- connection-based
  - ✱ connection must be established between sender and receiver first
  - ✱ connection is closed when communication is finished
- used for
  - ✱ reliable packet delivery
  - ✱ packet correctness
  - ✱ reliable order of packets
- work with TCP

## Types of Sockets (cont.)

### ◆ Datagram sockets

- connection-less socket
- used for
  - ✱ faster packet delivery
- work with UDP

### ◆ Raw sockets

- transport protocol is bypassed

## Socket Communications

### ◆ Steps in socket communications

- each endpoint opens a socket
- for stream sockets, a connection is established
- packets are sent and received
- each endpoint closes their socket

# Client-Server Model

- ◆ What is the client-server model?
  - a type of IPC architecture
- ◆ Characteristics
  - one server process receives requests and performs tasks
  - one or more client processes send requests to server

## Client-Server Model (cont.)

- ◆ Steps in establishing connection-based communications
  - server
    - \* create a stream **socket** on which to receive requests
    - \* **bind** the socket to its own IP address and port number
    - \* **listen** on the socket for incoming connection request from client
    - \* **accept** a connection request from client
    - \* receive (**recv**) data
    - \* **close** the socket
  - client
    - \* create a stream **socket** with which to connect to the server
    - \* **connect** to the server at its IP address and port number
    - \* **send** data
    - \* **close** the socket

## Client-Server Model (cont.)

### ◆ Steps in establishing connection-less communications

#### ● server

- \* create a datagram **socket** on which to receive requests
- \* **bind** the socket to its own IP address and port number
- \* **select** incoming request from client
- \* receive (**recvfrom**) data
- \* **close** the socket

#### ● client

- \* create a datagram **socket** with which to connect to the server
- \* send (**sendto**) data
- \* **close** the socket